VERJava: Vulnerable Version Identification for Java OSS with a Two-Stage Analysis

Abstract—The software version information affected by the CVEs (Common Vulnerabilities and Exposures) provided by the National Vulnerability Database (NVD) is not always accurate. This could seriously mislead the repair priority for software users, and greatly hinder the work of security researchers. Bao et al. improved the well-known Sliwerski-Zimmermann-Zeller (SZZ) algorithm for vulnerabilities (called V-SZZ) to precisely refine vulnerable software versions. But V-SZZ only focuses on those CVEs of which patches only have deleted lines.

In this study, we target Java Open Source Software (OSS) by virtue of its pervasiveness and ubiquitousness. Due to Java's object-oriented characteristic, a single security patch often involves modifications of multiple functions. Existing patch code similarity analysis does not give the patch existence result from the point of view of an entire patch, which would generate too many false positives for Java CVEs. In this work, we address these limitations by introducing a two-stage approach named VERJava, to systematically assess vulnerable versions for a target vulnerability in Java OSS. Specifically, vulnerable versions are calculated respectively at a function level and an entire patch level, then the results are synthesized to decide the final vulnerable versions. For evaluation, we manually annotated the vulnerable versions of 167 real CVEs from seven popular Java open source projects. The result shows that VERJava achieves the precision of 88% on average, significantly outperforming the state-of-the-art work V-SZZ. Furthermore, our study reveals some interesting findings that have not yet been discussed.

Index Terms—technological, patch analysis, vulnerability, Java OSS, vulnerable version identification, code similarity

I. INTRODUCTION

Nowadays, it is very common for software developers to use Open Source Software (OSS for short). According to the 2021 GitHub statistical report [1], more than 61 million new repositories were created in the last year. Besides, a recent study [2] shows OSSs usually contain a large number of vulnerabilities. To patch the vulnerabilities in OSSs promptly, maintainers often go to check the vulnerable versions of CVEs [3] given by NVD [4]. However, studies by Nguyen et al. [5], VIEM [6] and Bao et al. [7] show that CVE's version information in NVD is generally inaccurate. This could seriously mislead the repair priority for software maintainers, and greatly hinder the work of security researchers [8] [9] [10]. Therefore, it is of great importance to accurately identify the exact vulnerable versions affected by CVEs.

Identifying whether a software version is vulnerable to a known vulnerability in CVEs can be done by dynamically triggering the vulnerability. SemFuzz [11] and VULSCOPE [12] use fuzzing-based method for triggering. They first need a public PoC of CVE's reference version, and mutate certain data and control flows to generate triggerable PoCs for more versions. Dynamical approach introduces no false positives, however, it cannot ensure the vulnerability existence for the versions that are failed to construct a mutated PoC, which means that it may have false negatives. Moreover, this approach requires a runnable software and environment for each version under consideration, for the popular Java project Tomcat [13] which has hundreds of existing versions, the dynamic approach may suffer from scalability issues.

To address this problem in a relatively light-weight way, static analysis approach based on the blame feature of version-control systems also leads an active line of research (e.g., V-SZZ [7], SZZ [14], Alexopoulos et al. [15], OpenSZZ [16]). The Sliwerski-Zimmermann-Zeller (SZZ) algorithm is a well-known algorithm for identifying bug-inducing commits. A recent work V-SZZ [7] has improved it for identifying vulnerability-inducing commits. Essentially, V-SZZ uses the git blame command for the deleted lines in the patch to locate the earliest commit that introduce the deleted lines. The versions between the inducing commit and the fixing commit are considered to be the range of vulnerable versions. Such methods cannot handle the patches where the vulnerabilities are fixed by adding checks, namely, no deletion exists. Besides, this kind of algorithm heavily relies on the accuracy and completeness of commit IDs.

In response to these limitations, another instinctive reaction for identifying vulnerable version is to use the static approach based on the code similarity analysis between a target version and the patch code. There are a wealth of work on patch code analysis, while, aiming at the search for similar vulnerable codes, e.g., MVP [17], VCCFinder [18], ReDeBug [19] and VUDDY [20]. They focus on C and C++ projects and mainly identify the similar vulnerable code at the function level or at the patch-fragment (i.e., lines of consecutive modifications) level. In this work, we target Java OSS by virtue of its pervasiveness and ubiquitousness. According to our empirical analysis, security patches for Java OSS often involve multiple functions, specifically, of which portion is 61.7% in our collected dataset. Whether a version is vulnerable should be decided by considering the overall status of all patch functions, namely, at the entire patch level. Thus, existing patch analysis works would introduce many false positives by considering only fine-grained code unit. Moreover, the task of vulnerable version identification is different from the one of similar vulnerability searching. It is not clear whether some heavy static analysis methods like code normalization, forward/backward slicing often utilized in similar vulnerability searching are still needed in the task of vulnerable version

identification.

To address the above questions, in this paper, we first conduct a large-scale manual annotation and empirical analysis on Java patches and OSS versions, summarize the statistics with serveral observations. An important one is that thanks to Java's object-oriented characteristic, a modification in one method may bring according modifications in its sibling or interface classes, a single security patch often involves modifications of multiple functions. Under the guidance of the observations, we propose a light-weight and accurate approach "VERJava" for vulnerable version identification at the entire patch level. By "light-weight", we mean that no heavy compilation tool chain is involved to perform complex code slicing, thus VERJava can be easily deployed. By "accurate", we mean that the evaluation performed on the labelled dataset shows a rather satisfying precision of VERJava.

The key rationale behind our idea is that: since a Java patch often involve multiple functions, and software maintainers naturally apply an entire patch when fixing a vulnerability, if the majority of functions are similar to the post-patch code, the corresponding version is more likely to be patched. Specifically, we use a two-stage analysis. In the first stage, we perform a function level vulnerability existence analysis, namely, the versions where an involved function is un-patched are calculated using code similarity analysis. In the second stage, we perform an entire-patch level vulnerability existence analysis, namely, whether a target version is vulnerable is deduced from the proportion of patched functions.

To the best of our knowledge, this is the first work that considers the special security patch characteristic for the object-oriented language Java and identifies vulnerable versions at the entire patch level. The contributions of this paper are summarized as follows:

- For empirical analysis and evaluation, we manually annotated the vulnerable versions of 167 real CVEs from seven popular Java open source projects. A package of the detailed dataset is made publicly available[1].
- We propose a two-stage, i.e., function-level and patch-level, approach, named VERJava, to systematically assess vulnerable versions for a target vulnerability in Java OSS. Evaluation shows that the identification precision of VERJava achieves 88% on average, significantly outperforming the state-of-the-art work V-SZZ.
- Some interesting and valuable findings concerning the number of functions in a Java security patch and the object-oriented patch feature are summarized. They are helpful in inspiring the work in this paper, and hopefully be enlightening as well in works like similar vulnerability searching.

The remainder of this paper is organized as follows. In Section II, we use two examples to illustrate the limitation of existing approaches and motivate the idea of VERJava. Section III describes the data collection and manual vulnerability annotation efforts. Section IV outlines the design of

[1]https://anonymous.4open.science/r/java-dataset-FA7E/README.md

```
1  --- 7.0.x/.../authenticator/FormAuthenticator.java
       (revision 1408043)
2  @@ -404,6 +405,15 @@
3  ...
4  +    Session session = request.getSessionInternal(false);
5  +    if (session != null) {
6  +      Manager manager =
       request.getContext().getManager();
7  +      manager.changeSessionId(session);
8  +      request.changeSessionId(session.getId());
9  +    }
10 ...
```

Listing 1: The first motivating example from CVE-2013-2067

```
1  // commit 2235894210c75f624a3d0cd60bfb0434a20a18bf
2  --- /src/.../impl/SubTypeValidator.java
3  @@ -0,0 +1,98 @@
4  ...
5  +    if (full.startsWith(PREFIX_STRING)) {
6  +      for (Class<?> cls = raw; cls != Object.class;
       cls = cls.getSuperclass()) {
7  +        String name = cls.getSimpleName();
8  +        if ("AbstractPointcutAdvisor".equals(name)||
       "AbstractApplicationContext.equals".equals(name)){
9  +          break;
10 +        }
11 +      }
12 +    }
13 // SubTypeValidator.java in Jackson-databind version
       2.9.10
14     if (raw.isInterface()) {
15       ;
16     } else if (full.startsWith(PREFIX_SPRING)) {
17       for (Class<?> cls = raw; (cls != null) && (cls
            != Object.class); cls =cls.getSuperclass()){
18         String name = cls.getSimpleName();
19         if ("AbstractPointcutAdvisor".equals(name)||
             "AbstractApplicationContext".equals(name)){
20           break main_check;
21         }
22       }
23     } else if (full.startsWith(PREFIX_C3P0)) {
24       if (full.endsWith("DataSource")) {
25         break main_check;
26       }
27     }
```

Listing 2: The second motivating example from CVE-2017-17485

VERJava. Subsequently, Section V outlines the evaluation, while Section VI discusses the threat to validity and possible future improvement. Related work is discussed in Sect. VII and Sect. VIII concludes the paper.

II. CHALLENGES AND INSIGHTS

We choose the security patches of Tomcat's CVE-2013-2067 [21], Jackson-databind's CVE-2017-17485 [22] and Tomcat's CVE-2011-0013 [23] as examples to illustrate the limitations of existing methods, challenges, and our insights into addressing these challenges. Listing 1 shows the CVE-2013-2067 patch snippet for 7.0.x branch in *forwardToLoginPage()*. Listing 2 shows the patch snippets of CVE-2017-17485 in *validateSubType()* and the same function code in version 2.9.10. Listing 3 shows the patch of CVE-2011-0013 for branch 6.0.x.

```
1   --- 6.0.x/.../manager/HTMLManagerServlet.java (revision
        1057269)
2   @@ -407,10 +407,11 @@
3                 args = new Object[7];
4                 args[0] = URL_ENCODER.encode(displayPath);
5   -             args[1] = displayPath;
6   -             args[2] = context.getDisplayName();
7   -             if (args[2] == null) {
8   +             args[1] = RequestUtil.filter(displayPath);
9   +             if (context.getDisplayName() == null) {
10                    args[2] = " ";
11  +             } else {
12  +                 args[2] =
        RequestUtil.filter(context.getDisplayName());
13                }
14  --- 6.0.x/.../manager/StatusTransformer.java (revision
        1057269)
15  @@ -575,7 +575,7 @@
16  ...
17  -             writer.print(webModuleName);
18  +             writer.print(filter(webModuleName));
19  ...
20  @@ -650,7 +650,7 @@
21  ...
22  -         writer.print(name);
23  +         writer.print(filter(name));
24  ...
25  @@ -778,11 +778,11 @@
26  ...
27  -         writer.print(servletName);
28  +         writer.print(filter(servletName));
29  ...
30              for (int i = 0; i < mappings.length; i++) {
31  -             writer.print(mappings[i]);
32  +             writer.print(filter(mappings[i]));
33  ...
```

Listing 3: The third motivating example from CVE-2011-0013

A. Limitations of Existing Work

Vulnerability patches can be fixed in a variety of ways, such as removing or correcting the code that caused the vulnerability, or adding a checking code. We studied the patch codes of 167 real CVEs and found that 61 of them are fixed by adding a checking code, like the example in Listing 1. V-SZZ and other methods that based on the blame feature of version-control systems cannot handle such patches (marked as *Limitation 1)*. Our method considers multi-kind patches and can analyze and give the most accurate results for fixed patches by adding a checking code.

Secondly, due to the ubiquitous situation of multi-branch management in Java, different commit IDs are usually used when submitting commits. In order to accurately identify the vulnerability version on each branch, the method based on git blame needs to obtain the corresponding repair commit of the vulnerability on all branches, and the collection process of this information is more labor-intensive. In the analysis of seven popular Java OSS, it is found that only Tomcat maintains and repairs patches for each branch when publishing CVE, other projects basically only provide one patch, and it takes extra human resources to collect all branches' security fixes from issue history information (marked as *Limitation 2)*. For our method, we consider patch semantic information. We analyze according to the patch's content, and we can get accurate results with only one patch for the cases where the fixes are on different branches but the fixes are similar. In Listing 2,

we only got the patch for the 2.9.x branch of Struts, but we can detect that the vulnerability exists in other early branch simultaneously (like branch 2.6.x, 2.7.x, 2.8.x, etc.), and V-SZZ will cause false negatives. For the case where multiple branch patches can be fetched, we distinguish the patches of each branch, analyze each branch separately, and conduct a comprehensive evaluation of the results of each branch.

B. Challenges

We knew that CVE-2017-17485 was fixed in version 2.9.4 for branch 2.9.x. According to experience, there should be no CVE-2017-17485's vulnerability in versions 2.9.4 and later. With the evolution of the version and the code change, the code of *validateSubType()* may change. For example, the code of *validateSubType()* in version 2.9.10 shown in Listing 2. The conditional and loop statement has been modified, and a new conditional statement has been added, which is different from the patch code. Using simple added lines to consider the existence of the vulnerability will cause misjudging version 2.9.10 was vulnerable, so it is a challenge to consider the difference between target code and post-patch accurately (marked as *Challenge-I: code evolution processing)*? Here, we use the function similarity determination. We can comprehensively consider whether the current function is repaired through the similarity analysis between pre-patch and post-patch functions for the samples with code transitions (sees in IV-B).

For the example in Listing 3, CVE-2011-0013 not only fixes this one function but also fixes four function locations. The idea of patching is to verify the print function to prevent sensitive information from being seen by users. So all four functions are determined to be security fixes by manually annotating. After analyzing an extensive number of patches, we found that Java patches generally involve multiple function locations. How do we comprehensively estimate the version range of a vulnerability when the vulnerability exists in multiple functions in multiple files (marked as *Challenge-II: multi-location fix handling)*? We use the CVE Existing Vulnerability Version calculation method to handle the multi-location fix case (sees in IV-C).

III. DATASET CONSTRUCTION AND MANUAL ANNOTATION

To investigate the vulnerability fixing patterns of Java OSS, and propose an effective vulnerable version identification method, we first need a ground truth containing vulnerabilities and their affected software version lists. In this section, we describe the selection rules of target projects and CVEs (Sec. III-A), and then we manually annotate the vulnerable versions for each project and each CVE, the annotating principle is explained in Sec. III-B. Finally, two findings are given based on the statistics of labeled data (Sec. III-C).

A. Dataset Creation

The selection rules of target projects and known vulnerabilities to be manually annotated are presented below.

TABLE I: Number of CVEs, Commits, Functions and Versions we collected per project.

Target	Stars	CVEs	Commits	Functions	Versions
Tomcat	6k	45	125	430	144
Jackson-databind	3k	62	68	104	46
Struts	1.1k	14	18	114	41
Jenkins	18.7k	16	21	61	72
Liferay-portal	1.9k	20	99	200	13
Spring-security	6.7k	3	3	44	15
Spring-framework	47.1k	7	10	107	42
Total	-	167	344	1060	373

*1) **Target project selection criteria**:* To explore the characteristics of Java vulnerability patches, the dataset selection in our work is as follows: We first referred to the dataset of Java projects collected by Ponta et al. [24], which is also referred to in V-SZZ. In total, Ponta's dataset has 624 CVEs from more than 200 Java OSS. The TOP-6 projects w.r.t. the number of CVEs of which patches can be obtained are selected. Moreover, we searched on NVD details and selected Liferay-portal which has the most CVEs excluding the pre-selected TOP-6 projects. Due to the manual annotation overhead, we decided to first analyze the 7 projects, which are Tomcat [13], Jenkins [25], Struts [26], Jackson-databind [27], Spring-framework [28], Spring-security [29], Liferay portal [30].

*2) **CVE selection criteria**:* After the target projects have been determined, the selected CVEs are preferably to cover the following conditions as much as possible:

- The most important thing is that the patch file for the selected CVE can be directly obtained through NVD's CVE description page, or indirectly obtained through cross-searching NVD's CVE description page and the project's issue list;
- The selected CVEs need to cover a wide period of a project's lifetime;
- The selected CVEs need to cover different types of vulnerabilities (like path traversal, XSS, information disclosure, authentication bypass, code execution, DoS, etc.).
- The selected CVE needs to cover the various numbers of patch lines. There are cases where more than one hundred lines of code are modified in a Java vulnerability patch. There are also cases where only one line is modified in a vulnerability fixing patch(e.g., CVE-2008-0128's patch modifies only one line).

Eventually 167 CVEs are collected, and the details of the dataset are shown in Table I. The first and the second columns give the target OSS name and its number of github stars. The third column gives the total number of CVEs returned from a search in the NVD. The fourth column gives the corresponding sum of the commit numbers for those CVEs. Note that for the project Tomcat, the total commit number (i.e., 125) is nearly three times of the CVE number(i.e., 45), the reason is that Tomcat separately manages vulnerability fixing commits for each individual branch; for the project Liferay-portal, the total commit number (i.e., 99) is nearly 5 times of the CVE number

(i.e., 20), the reason is that this project tends to repeatedly submit patches for the same CVE. The fifth column shows the number of functions involved in those commits. And the last column displays the number of versions selected from these projects, which cover mainstream versions and also a wide range of the project's lifetime.

B. Manual Annotation Principles

To link each vulnerability with its affected versions, we perform static and manual auditing, without using dynamic ways to trigger it. Intuitively speaking, we think a software version is vulnerable if it contains the vulnerable code and does not contain the fixing code. Undoubtedly, objective annotation principles are essential for ensuring the accuracy of the subsequent experiment. We explain our manual annotation principles in details by considering the following three main situations.

*1) **Situation I: The function or file where the patch fragment is located does not exist in the target version's code**:* In that case, we think that the absence of patch function or file tends to show the absence of the vulnerability, thus annotate the target version as non-vulnerable. For example, all the functions involved in the fixing commit of CVE-2017-9805 do not exist before the version 2.1.0 of Struts. So we think before 2.1.0, Struts is not affected by CVE-2017-9805.

*2) **Situation II: The target code is the same as the pre-patch or the post-patch code**:* If the target code's counterpart is identical to the pre-patch code, it is annotated as vulnerable, and if the snippet is identical to the post-patch code, it is marked as fixed. For example, the pre-patch code for CVE-2020-35490 strictly matches the code for versions 2.9.5 to 2.10.1 of Jackson-databind. So we think the versions 2.9.5-2.10.1 of Jackson-databind are vulnerable.

*3) **Situation III: The target code is different from the pre-patch and the post-patch code**:* Differences include all kinds of changes in the target code, such as line missing, variable or method renaming, operator and control flow changes, etc. The example in Listing 2 displays that CVE-2017-17485's code in 2.9.10 is different from the patch code.

For this situation, we use the manual auditing expertise to determine whether the vulnerability is fixed in the current version through a comprehensive analysis of the vulnerability description, the patch code and the target code. The annotating principle in this situation is clarified in the following. We divide fixing patches into two categories, of which the first has both added and deleted lines, and the second has only one type of line modification. The added and deleted lines mentioned below are referred to as those critical code modifications identified by auditing expertise.

For the cases where a vulnerability fixing patch contains both added and deleted lines. We use two code-reviewing rules: (1) The existence of the added line in a target code takes precedence over the deleted lines. The empirical experience formed during manual labeling process shows that deleted lines may often appear in multiple places in the same function, while added lines recur less frequently. Namely, if some key

Fig. 1: Statistics of the proportion of the number of repaired functions in patches.

TABLE II: The data set we selected to analyze the actual proportion of annotation situations.

Target	CVE Num	Version Num
Tomcat	40	30
Jackson-databind	62	17
Struts	13	13
Jenkins	16	20
Liferay portal	19	13
Spring-security	3	10
Spring-framework	5	15
Total	**158**	**118**

TABLE III: The statistics of annotation situations.

Situation ID	Num	Percent	Merged Num	Percent
Situation I	1895	26.54%	1895	26.54%
Situation II (pre-patch)	2282	31.97%	4673	65.46%
Situation II (post-patch)	2391	33.49%		
Situation III (pre-patch)	60	1.84%	571	8.00%
Situation III (post-patch)	511	7.16%		

added lines identified by auditing expertise exist in the target code, the vulnerability is considered to be fixed in the target version, regardless of whether the deleted lines exist or not. (2) When the added lines do not exist in the target code but the deleted lines exist, it is assumed that there is a vulnerability in the target code.

For the cases where a vulnerability fixing patch contains only one type of line modification. They can be further divided into two categories: (1) There are no deleted lines in the patch (the vulnerability is fixed by adding checks, functions, classes, files, etc.). These cases can be analyzed only through the information given by the added lines. Generally speaking, if the added lines do not exist in the target code, the vulnerability is considered as unfixed in the target code, otherwise, the target code is considered to be secure. (2) There are no added lines in the patch. We can determine the target code is vulnerable or not only by the existence of the deleted lines. If the deleted lines exist in the target code, it is determined that there is a vulnerability in the target code, otherwise, the target code is considered to be secure.

C. Statistics and Findings

Finding 1: Vulnerability patches usually involve multiple functions in Java, of which percentage is about 61.7% in our dataset.

Figure 1 shows the percentage distribution of the number of functions involved in patches in each project. The red, yellow, light green and dark green colors represent the percentages of the cases where one function, 2-10, 11-20 and more than 20 functions being modified in one patch, respectively. The last row shows the overall distribution of the number of the fixing functions in the seven projects. According to the statistics shown in Figure 1, it indicates that Java vulnerability patches usually involve multiple function locations. Among the 167 CVEs, only 64 CVEs are repaired in one function, and the rest are multi-function repaired (61.7%). In particular, there are 10 CVEs with more than 20 repaired functions(6.6%).

Finding 2: The situation in which the code of each version is different from the pre-patch and post-patch code

occupies a small portion in Java OSS, of which percentage is only 8% in our labelled dataset.

To investigate the actual proportion of each annotation situation introduced in Sec. III-B, we perform a statistic analysis on a sub-dataset of which details are shown in Table II. The first column gives the target OSS name. The second column gives the total number of CVEs we selected (the CVEs whose numbers of patched functions are more than 20 are not counted). The third column gives the number of versions randomly selected at various stages of the target OSS's lifetime. In total, 158 CVEs and 118 versions are selected for labelling.

We compare each function involved in each patch against the counterpart function in each target version. For example, only one function is modified in the patch for CVE-2020-36184 of the project Jackson-databind. We will conduct seventeen comparisons for the seventeen versions we selected, to analyze which type of annotation situation each version belongs to. In total, we performed 7139 comparisons, and the detailed data are shown in Table III. The second column represents the statistics of the number of comparisons, and the third column represents the corresponding percentage. The fourth and fifth columns are the merged pre-patch and post-patch statistics. Through data analysis, we found that there are only 8% of the comparisons indicates that there is a discrepancy between the target version and the patch code (Situation III), the rest are the cases where either the patched files or functions are absent in the target version (26.54% for Situation I), or the patch code is identical to the target version (65.46% for Situation II).

The above two findings are helpful in guiding our algorithm design for vulnerable version identification. Since a Java patch often involves multiple-function modification (Finding 1), the designed algorithm should be aware of the fact that the analysis of whether a version is vulnerable should be conducted by comprehensively considering the patching status of all involved functions. Besides, since a majority of versions

is identical to the pre/post-patch code (Finding 2), we tend to not introduce immediately some heavy normalization or hash methods when designing the code matching algorithm.

IV. APPROACH AND DESIGN

Fig 2 shows the overview of VERJava, which mainly contains three steps. The information collection (Sec. IV-A) takes a target project's name and a CVE ID as input, uses a customized crawler to fetch the CVE's patch and collect multiple versions' source code for the target project. The function-level vulnerable version calculation step (Sec. IV-B) takes patch functions and multiple versions' source code as input and generates a list of vulnerable versions for each patch function. The patch-level vulnerable version calculation step (Sec. IV-C) takes all patch functions' vulnerable version lists as input and generates the vulnerable version list for the entire patch. Finally, we simply merge the results of all patches to get the final vulnerable versions of the CVE.

In Fig. 2, we illustrate the workflow of VERJava by using Tomcat's CVE-2011-0534 as an example. This CVE has a fixing patch with an issue management ID 1066313, containing five patch functions. At the information collection step, the patch functions and all versions' source code are collected. At the function-level vulnerable version calculation stage, a list of versions where each involved function is not patched are calculated respectively. For example, the involved function *InternalNioInputBuffer_expand()* is analyzed to be not patched in Tomcat versions 6.0.0-6.0.18, 7.0.0-7.0.6 and 8.0.8-8.0.x. Note that the vulnerable function searching is performed over the versions of all branches, that is why the result contains versions from different branches. At the patch-level vulnerable version calculation stage, basically, since almost all the five functions are un-patched in versions 6.0.0-6.0.18 and 7.0.0-7.0.6, these versions are considered vulnerable to CVE-2011-0534. Each step is further presented in details below.

A. Information Collection

We use different customized crawlers to collect CVE patches and multi-version source code of the target project. The CVE patch files (i.e., diff files) are obtained mainly in three ways: (1) There is a reference link on the CVE description page of NVD, which can directly bring us to the GitHub or SVN patch page. (2) According to the links on the CVE description page, we can find an internal management number of the project's issue list. Then we go to cross-search the project issue list according to the number to obtain the fixing commit. (3) We search for the CVE ID in the description of issue list of the target project and locate the corresponding issue and the fixing patch.

Furthermore, we split each obtained patch into functions, extract function information, such as file location, file name, function name, line numbers and line contents for added and deleted lines, etc, then store them in a JSON format.

B. Function-level Vulnerable Version Calculation

In this stage, we will decide the versions which are vulnerable at a function level, namely, the versions where an involved function is not patched. Preferably, if all deleted lines in a patch exist and all added lines in a patch do not exist in a function of a given version, then the function is vulnerable. However, patched functions may get changed as a project evolves. To tolerate the robustness, whether a function f of a given version V is vulnerable should be decided by synthesizing the degree of similarity between the target function and the patch code (to solve *Challenge-I: code evolution processing*). The decision algorithm is presented in Algorithm 1. We calculate the similarity between the target function and the pre/post-patch functions, and decide the vulnerability existence for the function under three different patch types: there are both added and deleted lines, only deleted lines and only added lines in the patch, respectively (Lines 6, 9 and 12 in Algorithm 1).

We denote the similarity between the target function and the pre-patch function by $delSim$, the similarity between the target function and the post-patch function by $addSim$. We use $TargetFunc$ to represent the function f body's lines in the version V's source code. The symbols $delLineSet$ and $addLineSet$ represent the sets of the deleted lines and the added lines of the patch function, respectively. For each line of source code, the tabs and spaces before and after it are deleted, and if a source code sentence spans multiple lines, we merge them into one line. The notation $\#$ is used to represent the number of items in a set. The formulas for calculating the function similarities are given below.

$$delSim \leftarrow \frac{\#(delLineSet \bigcap TargetFunc)}{\#delLineSet}$$

$$addSim \leftarrow \frac{\#(addLineSet \bigcap TargetFunc)}{\#addLineSet}$$

To compare the similarities in a robust way, we use two thresholds: the added line existence threshold (called $tAdd$) and the deleted line existence threshold (called $tDel$). When
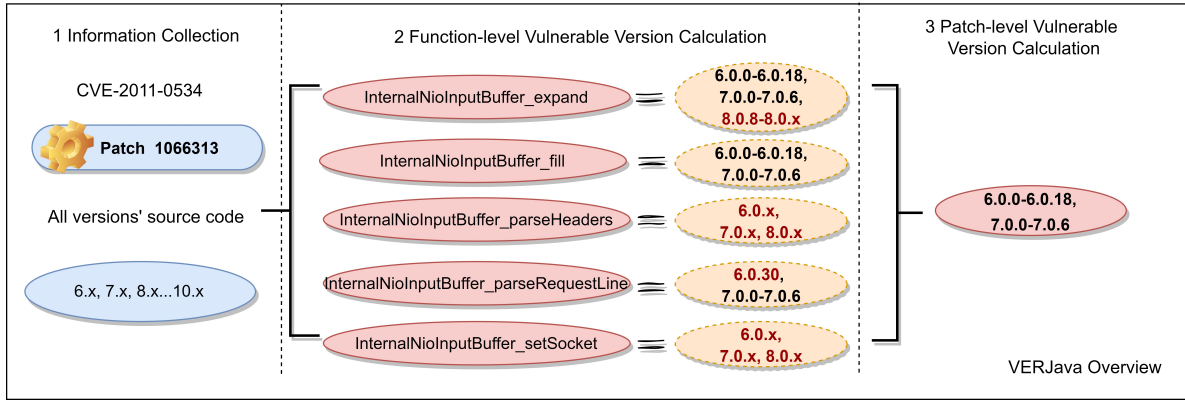

Algorithm 1 Vulnerable Function Calculation

1: Input: $TargetFunc$: function f's source code in a given version V, $Patchfunc$: function f's patch information in a given CVE
2: Output: whether $TargetFunc$ is vulnerable in V
3: $deleteLineSet \leftarrow$ the set of $Patchfunc$'s deleted lines
4: $addLineSet \leftarrow$ the set of $Patchfunc$'s added lines
5: $\#$ denotes the number of items in a set
6: **if** $addLineSet \neq \emptyset$ and $deleteLineSet \neq \emptyset$ **then**
7:     **if** $delSim \geq tDel$ and $addSim \leq tAdd$ **then**
8:         **return** vulnerable
9: **else if** $addLineSet == \emptyset$ **then**
10:     **if** $delSim \geq tDel$ **then**
11:         **return** vulnerable
12: **else if** $deleteLineSet == \emptyset$ **then**
13:     **if** $addSim \leq tAdd$ **then**
14:         **return** vulnerable
15: **return** safe

Fig. 2: Overview of VERJava's Approach

Algorithm 2 Patch-level Vulnerable Version Calculation

1: Input: $AllPFun$: the set of functions in a given patch P, $VersionList$: the list of target software versions
2: Output: vulnerable versions w.r.t. P
3: $TotalNum \leftarrow$ the number of functions in $AllPFun$
4: $VulNum \leftarrow$ the number of vulnerable functions w.r.t. P in version V
5: Initialize $vulnerableVersions$ to be empty.
6: **for** each $V \in VersionList$ **do**
7:     **for** each $f \in AllPFun$ **do**
8:         **if** f does not exist in V **then**
9:             $TotalNum$ -= 1
10:     **if** $(TotalNum > 3 \bigwedge VulNum\,/\,TotalNum \geq T) \vee (TotalNum \leq 3 \bigwedge VulNum\,/\,TotalNum == 1)$ **then**
11:         $vulnerableVersions$.insert(V)
12: **return** $vulnerableVersions$

$delSim$ is higher than $tDel$, and $addSim$ is lower than $tAdd$, we consider $TargetFunc$ as vulnerable (Line 6-8). For patches with only one type of modification, we only need to consider the similarity with the pre-patch (Line 9-11) or post-patch (Line 12-14) function. Experiments on the threshold sensitivity have been conducted in Sec. V and the results show that the best effect is achieved when $tdel$ is 1 and $tAdd$ is 0.9. Therefore, $tDel$ and $tAdd$ are empirically set to be 1 and 0.9, respectively.

C. *Patch-level Vulnerable Version Calculation*

Theoretically, when a patch is applied, all involved functions are fixed, so the vulnerable versions of all patch functions should be the same. However, due to code modifications in the version evolution process, the vulnerable versions calculated by the previous stage are different for different functions in the same patch, as already seen in Fig. 2. Considering the prevalence of multi-function fixing in Java patches (to solve *Challenge-II: multi-location fix handling*), we need to deal with all patch functions' results in a comprehensive way. The decision algorithm is stated in Algorithm 2.

$TotalNum$ represents the number of functions in a given patch P. First we count how many patch functions are present in a given version V. There are code refactoring in some versions as a project evolves. If a patch function does not exist in V any more, we subtract 1 from $TotalNum$ (Lines 7-9 in Algorithm 2).

Ideally, the version V is considered vulnerable with respect to (w.r.t.) P when all vulnerability version lists of functions in P contain V. However, this condition is so strict that it may introduce false positives, as there are cases where patched lines are modified as the code evolves. Therefore, here we use a threshold T (Line 10) to control false positives. It is empirically to be 0.8 based on threshold sensitivity analysis in Sec. V. But for the case where the number of patch functions, i.e., $TotalNum$, is less than or equal to three, the ratio $VulNum/TotalNum$ varies greatly and the threshold is strictly required to be 1 (Line 10-11).

Software maintainers usually manage Java projects in multiple branches (e.g., Tomcat maintains multiple branches such as 10.0.x, 9.0.x, 8.0.x, 8.5.x, 7.0.x, etc.). Therefore, when fixing a CVE, they also submit patches separately for different branches. Thus we separately analyze the vulnerable versions w.r.t. each patch from different branches, merge the results of all patches and obtain a final list of versions vulnerable to the CVE.

V. EVALUATION

In this section, we evaluate VERJava on 167 Java real-world CVEs (detail shows in Table I). To the best of our knowledge, only tools in the SZZ series directly target vulnerable OSS version identification and use also static analysis. V-SZZ is the state-of-the-art one among them, so only it was selected for comparison. We evaluate VERJava regarding accuracy and time overhead in identifying the CVE version range.

A. *Effectiveness of VERJava*

Our dataset contains seven real-world open-source Java projects, including 167 CVEs, and extracts 344 patches from them. Our method takes patch function as the unit and collects 1060 patch functions.

TABLE IV: Accuracy (i.e., True Positive, False Positive and False Negative) of V-SZZ, ΔV-SZZ and VERJava. Prec. stands for precision(True Positive Number / Ground Truth Number). T-FN/T-Prec are calculated using the same set as VERJava (including those patches that V-SZZ cannot handle), while ordinary FN/Prec is obtained by excluding those cases.

| Tool	GT	V-SZZ						ΔV-SZZ						VERJava			
		TP	FP	FN	Prec.	T-FN	T-Prec.	TP	FP	FN	Prec.	T-FN	T-Prec.	TP	FP	FN	Prec.
Tomcat	45	9	1	0	90%	35	20.0%	33	1	4	86.8%	21	73.3%	39	6	0	86.7%
Jackson-databind	62	3	2	2	42.9%	57	4.8%	3	2	2	42.9%	57	4.8%	57	5	0	91.9%
Struts	14	5	0	4	55.6%	9	35.7%	6	0	8	42.9%	8	42.9%	12	0	2	85.7%
Jenkins	16	7	5	0	41.7%	4	43.8%	7	6	0	53.8%	3	43.8%	12	1	3	75%
Liferay portal	20	-	-	-	-	-	-	-	-	-	-	-	-	17	3	0	85%
Spring-security	3	0	0	2	0%	3	0%	0	1	2	0%	2	0%	3	0	0	100%
Spring-framework	7	0	2	0	0%	5	0%	1	4	0	20%	2	14.3%	7	0	0	100%
Total	167	24	10	8	57.1%	133	14.4%	50	14	16	62.5%	103	29.9%	147	15	5	88.0%

Fig. 3: Version-level tool precision results.

Overall Results. Table IV shows the results of VERJava in identifying the Java vulnerable version range. Among the 167 CVEs, VERJava can correctly identify the version of 147 CVEs (88% true positive). And VERJava has 9% false positives (FP) and 3% false negatives (FN). Mention that only all the versions marked correctly will be determined to be true positive (TP). For example, the CVE-2017-1000391 annotation should report 61 versions, and VERJava accurately reports 61 versions that exactly match the version list of the annotation so that the results of CVE-2017-100391 is marked as a TP. Similarly, CVE-2009-3742 reported one more version, but we will mark it as a FP. The used calculation method for FP/FN is rigorous. We add the measurement of the average percentage of matched versions over all vulnerable versions. The results show in Figure 3. In terms of the precisions, VERJava achieves the best performance. Therefore, we believe that VERJava has a good accuracy in identifying vulnerable version ranges.

False Positive and False Negetive Analysis for VERJava. VERJava still has 9% FP. In the example CVE-2013-2067 in Listing 1, there are only added lines in the patch for version 7.0.33. But in version 7.0.100, the added lines have class and method modifications. We confirmed through manual auditing that the 7.0.100 version has fixed the vulnerability. While VERJava's method cannot determine if 7.0.100 has been fixed because of the significant semantic and code changes, which leads to FPs. There are five FNs in VERJava, all of which are since some versions only exist in one or a few patch functions' results, which account for a small proportion of the total functions. If one patch function in the target version is

vulnerable, we determined that this version has a vulnerability during manual annotation. VERJava considers according to the proportion of patch functions. In addition, the versions that were found to be FN are early versions, and more attention is paid to the version with a relatively recent release time in the actual vulnerability analysis, so the impact of under-reporting in practical applications is small.

B. Comparison with V-SZZ

We found that most CVEs could not be successfully run on V-SZZ (close to 70%) for the reasons: (1) The patches fixes only by added lines. In Java patch fixes, such cases account for a large proportion (61 out of 167 CVEs). V-SZZ cannot handle patches fixes by added lines because of the limitation of the algorithm based on the blame, while VERJava can, which means VERJava can cover more patch types. (2) V-SZZ only processes patches with less than five lines deleted due to time consumption. And according to our statistics, many line number patch fixes are prevalent. Here we improve V-SZZ (asΔV-SZZ) to remove this limit and evaluate them separately.

Specifically, for multi-branch management, different commit IDs would be assigned to each submission . The method of V-SZZ highly dependent on the commit IDs. For example, CVE-2020-36183 submitted two fixing commits on the 2.0 and 2.9 branches. V-SZZ needs two different fixing commit IDs to determine the version range accurately, missing enough commit IDs could cause FN. Otherwise, VERJava uses patch content for analysis, which means we can accurately identify the vulnerability version range of CVE-2020-36183 based on only one commit. Moreover, several V-SZZ mistake reports are due to defects in the official commit management. There are situations such as wrong-submission of some tags or multiple submissions when submitting a commit (e.g., when fixing CVE-2020-11111, the software maintainer mistakenly submitted a patch, and then submitted a patch for revision). Liferay portal cannot use the git blame command because the file path is too long, so we did not use V-SZZ for this project which do not have results in Table IV's. The results show that VERJava has higher recognition accuracy and precision.

Time Overhead. Both VERJava and V-SZZ require a manual or semi-automated part of the patch collection. We first collect as many source code versions as possible, but V-

TABLE V: Performance Overhead of VERJava, V-SZZ, ΔV-SZZ

TOOL	CVE	MAX TIME	MIN TIME	AVERAGE TIME
VERJava	167	0.31min	10ms	0.71s
V-SZZ	42	2.65min	30ms	7.60s
ΔV-SZZ	80	28.69min	30ms	29.39s

SZZ requires git checkout operations for each analysis, so our analysis time is shorter in terms of analysis time. According to Table V, VERJava only needs 118.5s to analyze 167 CVEs, and the average analysis time of each CVE is 0.71s. The analysis time is positively correlated with the number of patch functions of CVEs. There is one patch function with the least time consuming which is only 10ms. The most time-consuming is CVE-2012-0022, which has 46 patch functions with 18.76s' time consumption. Moreover, the average time for V-SZZ to analyze each CVE is 30.22s which is worse than VERJava. The time consumption of ΔV-SZZ is worse than that of V-SZZ. After removing the five-element limit of V-SZZ, although more CVEs can be processed, the time consumption is greatly improved.

C. Threshold Sensitivity Analysis

Three thresholds (i.e., tAdd, tDel, T) are configurable in the steps of VERJava. The default configuration is 0.9, 1, 0.8, which is used in the experiment in Algorithm 1 and Algorithm 2. To evaluate the sensitivity of these thresholds to the accuracy of VERJava, we reconfigured one threshold and fixed the other two, and ran VERJava against the seven targets. As $tAdd$ and $tDel$ are used to determine whether a function is vulnerable, they were reconfigured from 0.1 to 1.0 by a step of 0.1. As T is adopted to determine whether CVE exists in version V, it was reconfigured from 0.1 to 1 by a step of 0.1. In total, 29 (i.e., 3 × 10 - 1) configurations of VERJava were run.

Fig. 4 presents the impact of three thresholds on precision, respectively. Overall, before $tAdd$ increased to 0.9, the precision was almost stable in most target systems. As $tAdd$ increased from 0.9 to 1.0, the precision decreased, since many TP were missed due to vulnerability patching method modification in version evolution. As $tDel$ increased from 0.1 to 1, the precision increased. Thus, we believe that 0.9 and 1 are good values for $tAdd$ and $tDel$, respectively. On the other hand, as T increased from 0.0 to 0.8, the precision increased. Thus, we believe 0.8 is a good value for T.

D. Interesting Findings

During our analysis results we found some other interesting findings.

Accuracy of NVD version labeling in seven target projects. We manually annotated our selected 167 CVEs. It takes about 20 minutes to mark the existence of a vulnerability per function. The annotation results show that NVD has incorrect version annotations for 121 of the 167 CVEs(72%). As shown in Figure 5, we count the accuracy of NVD version annotations in seven projects.

Sibling classes and interface class fixes. During the manual analysis of the patch, we found repairs to sibling classes and interface classes. For example, when CVE-2013-4322 was fixed, the same code was fixed in four classes. *AbstractHttp11Processor* and *Http11Processor* are interface classes, while *Http11AprProcessor* and *Http11NioProcessor* are siblings. According to statistics, 34 of the 167 CVEs have sibling repairs, and 9 have interface repairs. It means that a patch may involve the repair of multiple vulnerabilities. We can separate situations such as sibling and interface classes to represent multiple vulnerabilities. From an analysis point of view, it can enable finer-grained patch analysis and detection of similarity vulnerabilities. This finding may guide the technical route for object-oriented language vulnerability discovery.

VI. DISCUSSION

In this section, we discuss the threat to validity and possible future improvement of VERJava.

VERJava does not differentiate between secure and non-secure patches. We think that all the obtained patches are security fixes. Patches submitted have security fixes and non-security fixes. When doing manual annotation and tool analysis, we assumed that all modified codes in the patch are security-related. In the future, more great vulnerability detection results and potentially vulnerable classes and methods are acquired for security personnel to analyse conveniently.

VERJava does not consider the extract function. Vulnerabilities may exist, but their locations may be scattered among different features of earlier versions of the project. VERJava assume that extract function do not exist. We would build a function migration mapping relationship for the code refactoring based on version control system, and expand the search scope of vulnerable functions based on it, to mitigate the vulnerable function/file absence imposed by refactoring to some extent.

The vulnerable versions found by VERJava have not been verified. Like other static analysis methods, VERJava does not necessarily contain vulnerabilities because there is no dynamic vulnerability triggering. In addition, there are some vulnerabilities that require a particular setting to be enabled in the configuration file. Our work does not consider such cases. In the future, we would like to perform dynamic vulnerability trigger research on the results of static analysis or add analysis to configuration files during vulnerability analysis. The results of comprehensive configuration items and code audits are reported in a more accurate version range of vulnerabilities.

VII. RELATED WORK

We review the most closely related work in vulnerable version identification and patch analysis.

A. Vulnerable Version Identification

Sliwerski et al. [14] proposed a technique called Sliwerski-Zimmermann-Zeller (SZZ for short) that automatically locates fix-inducing changes. This technique first locates changes for bug fixes in the commit log, then tries to find the commit

Fig. 4: Precision vs. Three Threshold (x-axis denotes the value of threshold, and y-axis denotes precision)


Fig. 5: Software labelling accuracy statistics

which the modified line added called inducing commit. Then SZZ confirms a bug in the range between the inducing commit and the fixing commit. Many researchers have improved the SZZ method and enhanced tools (VSZZ [7], OpenSZZ [16], AG-SZZ [31], MA-SZZ [32], etc.). They are all tools for identifying the version range where a vulnerability exists. However, these methods based on the blame feature of version-control systems, can not handle patches which were fixed by adding code. The work of Nguyen et al. [5] considers both the added and deleted lines in the patch. It can perform relatively accurate vulnerability version range identification for Chrome and Firefox. However, this work does not differentiate between the characteristics of the patches. In addition, the current version identification works do not consider the characteristics of multi-branch management, multi-function repair, and multi-line repair of the object-oriented language's patch.

B. Patch Analysis

The target software of existing work is divided into analysis closed source software [33] [34] [35] and open source software [17] [18] [19] [20] [36] [37] [38] [39] .

The design of FIBER [33] insufficiently considers the code variance incurred by third-party customization and non-standard building configurations, which consequently limits the effectiveness of FIBER in practice. PDIFF [35] and BSCOUT [34] and made improvements for this weakness. But for similarity computation between binaries, PDIFF generates relatively heavy-weight code feature, i.e., sets of patch-related paths, which is less influenced by compilation and building configurations. While for OSS, we uses lightweight but effective code features. After mapping source codes to Java bytecodes, BSCOUT adopts line-level similarity analysis but does not consider from function-level perspective and neither Challenge I nor II is proposed in the paper.

Patch analysis work for open source software is mainly used for vulnerability identification. ReDeBug [19] is a well-known tool for patch presence analysis. Li et al. [37] [38] [39] made some advances in patch presence analysis. VulPacker [37] defines code-similarity algorithms to automatically detect patch. On this basis, they added deep learning algorithms (Vulpecker [38], VulDeePecker [39]), and the new methods of research achieved better efficiency. Another study of patch analysis is used to detect similarity vulnerabilities like MVP [17], VCCFinder [18], VUDDY [20]. They all focus on patch analysis in C/C++, focus on fixes within functions, and do not focus on the difference between patches in object-oriented and typical languages. Other studies exist for patch analysis. SecureSync [36] uses patches to analyze vulnerabilities caused by code cloning. Chen et al. [40] [41] provide solutions that can identify vulnerabilities and automatically patch vulnerabilities. Tian et al. [42] is used to evaluate patch correctness in predictor fixes. Although these works are different from our goals, some analysis and processing methods can be used for us to reference in patch analysis.

VIII. CONCLUSIONS

In this paper, we collected 167 CVEs across seven popular Java open-source software. We found that the version labels of 121 CVEs were inaccurate through analysis. Due to the object-oriented characteristics of Java patches, it has the characteristics of general multi-function repair, a large number of repair code lines, and multi-branch management. To handle these characteristics, we have proposed and implemented a novel approach named VERJava to identify the vulnerability version range. Our evaluation results have demonstrated that VERJava can significantly outperform the state-of-the-art approaches.

REFERENCES

[1] GitHub, "GitHub 2021 Report," https://octoverse.github.com/.
[2] M. Backes, S. Bugiel, and E. Derr, "Reliable third-party library detection in android and its security applications," in *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, 2016, pp. 356–367.
[3] CVE, "Common vulnerabilities and exposures," https://cve.mitre.org/.
[4] NVD, "National vulnerability database," https://nvd.nist.gov/.
[5] V. H. Nguyen, S. Dashevskyi, and F. Massacci, "An automatic method for assessing the versions affected by a vulnerability," *Empirical Software Engineering*, vol. 21, no. 6, pp. 2268–2297, 2016.
[6] Y. Dong, W. Guo, Y. Chen, X. Xing, Y. Zhang, and G. Wang, "Towards the detection of inconsistencies in public security vulnerability reports," in *28th USENIX Security Symposium (USENIX Security 19)*. Santa Clara, CA: USENIX Association, Aug. 2019, pp. 869–885.
[7] L. Bao, X. Xia, A. E.Hassan, and X. Yang, "SZZ for vulnerability: Automatic identification of version ranges affected by cve vulnerabilities." ICSE, 2022.
[8] O. Chaparro, J. Lu, F. Zampetti, L. Moreno, M. Di Penta, A. Marcus, G. Bavota, and V. Ng, "Detecting missing information in bug descriptions," in *Proceedings of the 2017 11th Joint Meeting on Foundations of Software Engineering*, 2017, pp. 396–407.
[9] Y. Dong, W. Guo, Y. Chen, X. Xing, Y. Zhang, and G. Wang, "Towards the detection of inconsistencies in public security vulnerability reports," in *28th USENIX Security Symposium (USENIX Security 19)*, 2019, pp. 869–885.
[10] X. Tan, Y. Zhang, C. Mi, J. Cao, K. Sun, Y. Lin, and M. Yang, "Locating the security patches for disclosed oss vulnerabilities with vulnerability-commit correlation ranking," in *Proceedings of the 2021 ACM SIGSAC Conference on Computer and Communications Security*, 2021, pp. 3282–3299.
[11] W. You, P. Zong, K. Chen, X. Wang, X. Liao, P. Bian, and B. Liang, "SemFuzz: Semantics-Based Automatic Generation of Proof-of-Concept Exploits," in *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS '17. New York, NY, USA: Association for Computing Machinery, 2017, p. 2139–2154.
[12] J. Dai, Y. Zhang, H. Xu, H. Lyu, Z. Wu, X. Xing, and M. Yang, "Facilitating vulnerability assessment through poc migration," in *Proceedings of the 2021 ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS '21. New York, NY, USA: Association for Computing Machinery, 2021, p. 3300–3317.
[13] Apache, "Tomcat," https://tomcat.apache.org/index.html.
[14] J. Śliwerski, T. Zimmermann, and A. Zeller, "When do changes induce fixes?" *ACM sigsoft software engineering notes*, vol. 30, no. 4, pp. 1–5, 2005.
[15] N. Alexopoulos, M. Brack, J. P. Wagner, T. Grube, and M. Mühlhäuser, "How long do vulnerabilities live in the code? a large-scale empirical measurement study on foss vulnerability lifetimes."
[16] V. Lenarduzzi, F. Palomba, D. Taibi, and D. A. Tamburri, "OpenSZZ: A free, open-source, web-accessible implementation of the szz algorithm," in *Proceedings of the 28th international conference on program comprehension*, 2020, pp. 446–450.
[17] Y. Xiao, B. Chen, C. Yu, Z. Xu, Z. Yuan, F. Li, B. Liu, Y. Liu, W. Huo, W. Zou *et al.*, "MVP: Detecting vulnerabilities using Patch-Enhanced vulnerability signatures," in *29th USENIX Security Symposium (USENIX Security 20)*, 2020, pp. 1165–1182.
[18] H. Perl, S. Dechand, M. Smith, D. Arp, F. Yamaguchi, K. Rieck, S. Fahl, and Y. Acar, "VCCFinder: Finding potential vulnerabilities in open-source projects to assist code audits," in *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security*, 2015, pp. 426–437.
[19] J. Jang, A. Agrawal, and D. Brumley, "ReDeBug: finding unpatched code clones in entire os distributions," in *2012 IEEE Symposium on Security and Privacy*. IEEE, 2012, pp. 48–62.
[20] S. Kim, S. Woo, H. Lee, and H. Oh, "VUDDY: A scalable approach for vulnerable code clone discovery," in *2017 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2017, pp. 595–614.
[21] Apache, "CVE-2013-2067 7.x branch commit," https://svn.apache.org/viewvc?view=revision&revision=1408044.
[22] FasterXML, "CVE-2017-17485 commit," https://github.com/FasterXML/jackson-databind/commit/2235894210c75f624a3d0cd60bfb0434a20a18bf.
[23] Apache, "CVE-2011-0013 6.x branch commit," https://svn.apache.org/viewvc?view=revision&revision=1057270.
[24] S. E. Ponta, H. Plate, A. Sabetta, M. Bezzi, and C. Dangremont, "A manually-curated dataset of fixes to vulnerabilities of open-source software," in *Proceedings of the 16th International Conference on Mining Software Repositories*, ser. MSR '19. IEEE Press, 2019, p. 383–387.
[25] Jenkins, "Jenkins," https://www.jenkins.io/.
[26] Apache, "Struts," https://struts.apache.org/.
[27] FastXML, "Jackson-databind," https://github.com/FasterXML/jackson-databind/wiki.
[28] Spring, "Spring-framework," https://spring.io/projects/spring-framework.
[29] ——, "Spring-security," https://spring.io/projects/spring-security.
[30] Liferay, "Liferay-portal," https://www.liferay.com/zh/locations.
[31] S. Kim, T. Zimmermann, K. Pan, and E. J. Jr. Whitehead, "Automatic identification of bug-introducing changes," in *21st IEEE/ACM International Conference on Automated Software Engineering (ASE'06)*, 2006, pp. 81–90.
[32] D. A. da Costa, S. McIntosh, W. Shang, U. Kulesza, R. Coelho, and A. E. Hassan, "A framework for evaluating the results of the szz approach for identifying bug-introducing changes," *IEEE Transactions on Software Engineering*, vol. 43, no. 7, pp. 641–657, 2017.
[33] H. Zhang and Z. Qian, "Precise and accurate patch presence test for binaries," in *27th USENIX Security Symposium (USENIX Security 18)*, Baltimore, MD, Aug. 2018, pp. 887–902.
[34] J. Dai, Y. Zhang, Z. Jiang, Y. Zhou, J. Chen, X. Xing, X. Zhang, X. Tan, M. Yang, and Z. Yang, "BScout: Direct whole patch presence test for java executables," in *29th USENIX Security Symposium (USENIX Security 20)*. USENIX Association, Aug. 2020, pp. 1147–1164.
[35] Z. Jiang, Y. Zhang, J. Xu, Q. Wen, Z. Wang, X. Zhang, X. Xing, M. Yang, and Z. Yang, "PDIFF: Semantic-based patch presence testing for downstream kernels," in *Proceedings of the 2020 ACM SIGSAC Conference on Computer and Communications Security*, 2020, pp. 1149–1163.
[36] N. H. Pham, T. T. Nguyen, H. A. Nguyen, and T. N. Nguyen, "Detection of recurring software vulnerabilities," in *Proceedings of the IEEE/ACM international conference on Automated software engineering*, 2010, pp. 447–456.
[37] Z. Li, D. Zou, S. Xu, H. Jin, H. Qi, and J. Hu, "VulPecker: an automated vulnerability detection system based on code similarity analysis," in *Proceedings of the 32nd Annual Conference on Computer Security Applications*, 2016, pp. 201–213.
[38] Z. Li, D. Zou, S. Xu, X. Ou, H. Jin, S. Wang, Z. Deng, and Y. Zhong, "VulDeePecker: A deep learning-based system for vulnerability detection," *arXiv preprint arXiv:1801.01681*, 2018.
[39] D. Zou, S. Wang, S. Xu, Z. Li, and H. Jin, "μVulDeePecker: A deep learning-based system for multiclass vulnerability detection," *IEEE Transactions on Dependable and Secure Computing*, vol. 18, no. 5, pp. 2224–2236, 2019.
[40] Y. Chen, A. E. Santosa, A. Sharma, and D. Lo, "Automated identification of libraries from vulnerability data," in *2020 IEEE/ACM 42nd International Conference on Software Engineering: Software Engineering in Practice (ICSE-SEIP)*, 2020, pp. 90–99.
[41] Y. Chen, A. E. Santosa, A. M. Yi, A. Sharma, A. Sharma, and D. Lo, *A Machine Learning Approach for Vulnerability Curation*. New York, NY, USA: Association for Computing Machinery, 2020, p. 32–42.
[42] H. Tian, K. Liu, A. K. Kaboré, A. Koyuncu, L. Li, J. Klein, and T. F. Bissyandé, "Evaluating representation learning of code changes for predicting patch correctness in program repair," in *Proceedings of the 35th IEEE/ACM International Conference on Automated Software Engineering*. New York, NY, USA: Association for Computing Machinery, 2020, p. 981–992.